

---

## Ansible Dynamic Inventory on AWS

---

### 1. Introduction to Ansible Inventory

Ansible needs an **inventory** to know *which servers to manage*.

#### Static Inventory Example

```
[web]
10.0.0.10
10.0.0.11
```

Static inventory works for small, fixed environments.

But in cloud environments like AWS, it becomes a problem because:

- New EC2 instances get created dynamically
- IPs change frequently
- Auto-scaling launches and terminates servers
- Tags define roles (web, devops, db)
- Manual updates are not feasible

---

This is where **Dynamic Inventory** becomes essential.

---

## 2. What is Dynamic Inventory?

Dynamic Inventory means Ansible **automatically discovers servers** by querying AWS APIs.

Instead of writing IPs manually, Ansible uses:

- AWS tags
- AWS metadata
- AWS regions

Ansible fetches:

- Public IP
- Private IP
- Instance ID
- Tags (Name, Role, Application, etc.)
- State (running/stopped)

This enables **fully automated, scalable, cloud-native deployments**.

---

## 3. Use-Case Example

You created an EC2 instance with:

- Tag: Name = centos

- Tag: `Role = devops`

You want Ansible to:

- Automatically detect this instance
- Group it under `role_devops` and `name_centos`
- Connect to it over SSH
- Run playbooks on it

All **without manually updating any inventory file.**

---

#### 4. Prerequisites

**On Ansible Master (Ubuntu Server):**

- Python 3
- Virtual environment support
- `boto3` & `botocore`
- Ansible AWS collection
- AWS CLI credentials
- SSH private key (PEM file)

Your SSH key is:

~/worklab.pem

---

## 5. Installing Requirements (Ubuntu 24.04 & 22.04)

Ubuntu now blocks global pip installs (PEP 668).

So we use a virtual environment.

---

### Step 1 — Install python venv

```
sudo apt update
```

```
sudo apt install python3.12-venv -y
```

---

### Step 2 — Create Virtual Environment

```
python3 -m venv myenv
```

```
source myenv/bin/activate
```

---

### Step 3 — Install boto3 & botocore

```
pip install boto3 botocore
```

---

### Step 4 — Install AWS Collection for Ansible

```
ansible-galaxy collection install amazon.aws
```

---

### Step 5 — Configure AWS credentials

```
aws configure
```

Provide:

- AWS Access Key
  - AWS Secret Key
  - Default region (e.g., `us-east-2`)
  - Output format json
- 

## 6. Creating the Dynamic Inventory File

Create:

```
/etc/ansible/aws_ec2.yml
```

Paste this **correct, production-ready configuration:**

```
plugin: amazon.aws.aws_ec2
```

```
regions:
```

```
- us-east-2
```

```
filters:
```

```
tag:Name: centos
```

```
tag:Role: devops
```

```
hostnames:
```

```
- ip-address # Correct field for public IP
```

```
keyed_groups:  
  - key: tags.Role  
    prefix: role  
  - key: tags.Name  
    prefix: name
```

### ✓ Why ip-address?

Because this is the correct AWS metadata field the plugin uses to fetch public IPs.

---

## 7. Validate Dynamic Inventory

Run:

```
ansible-inventory -i /etc/ansible/aws_ec2.yml --list
```

You should now see:

```
"hosts": ["18.218.218.228"]
```

This means dynamic inventory is working properly.

---

## 8. Fixing SSH Authentication

Initially, the error appeared:

```
no such identity: /home/ubuntu/.ssh/worklab.pem
```

Meaning Ansible was searching in the wrong folder.

Your key is located at:

~/worklab.pem

So use:

--private-key ~/worklab.pem

---

## 9. Determine the Correct SSH Username

Your AMI ID:

ami-028bd5cea26d40561

This AMI = **Ubuntu 22.04 LTS**

Therefore, the correct user is:

ubuntu

Not **centos**.

---

## 10. Test SSH Connectivity Using Ansible

Final working ping command:

```
ansible -i /etc/ansible/aws_ec2.yml all -m ping \  
-u ubuntu --private-key ~/worklab.pem
```

Successful output:

```
18.218.218.228 | SUCCESS => {
    "changed": false,
    "ping": "pong"
}
```

Now your dynamic inventory is verified.

---

## 11. Running a Playbook Using Dynamic Inventory

Create a file:

```
deploy.yml
```

```
- hosts: all
become: yes
tasks:
  - name: Install Git
    apt:
      name: git
      state: present
```

Run it:

```
ansible-playbook -i /etc/ansible/aws_ec2.yml deploy.yml \
-u ubuntu --private-key ~/worklab.pem
```

This installs git on the AWS instance automatically discovered via dynamic inventory.

---

## 12. Understanding AWS Dynamic Inventory Internals

When you run Ansible:

1. Plugin connects to AWS using boto3
  2. AWS returns all EC2 instances
  3. Ansible filters only:
    - o Name=centos
    - o Role=devops
  4. Ansible extracts the IP using:
    - o ip-address
  5. Groups servers into:
    - o role\_devops
    - o name\_centos
  6. Runs modules/playbooks on discovered instances
  7. No manual host updates required
- 

### 13. Best Practices

- ✓ Use IAM roles for EC2 Ansible Masters
- ✓ Never hardcode AWS credentials in playbooks
- ✓ Secure your PEM file:

```
chmod 400 ~/worklab.pem
```

✓ Use tagging strategy:

Environment = dev

Application = web

Role = backend

✓ Use separate dynamic inventories per environment:

- aws\_dev.yml
- aws\_stage.yml
- aws\_prod.yml

✓ Always test with:

```
ansible-inventory --graph -i aws_ec2.yml
```

---

## 14. Common Errors & Fixes

Error	Cause	Fix
public_ip host shown	Wrong hostname field	Use ip-address
Permission denied	Wrong user	Use ubuntu
key not found	Wrong path	Use ~/worklab.pem
PEP 668 pip error	Ubuntu restriction	Use python3 -m venv

Inventory empty	Wrong tags/region	Verify tags in AWS console
unreachable	SG restriction	Allow SSH from Ansible master IP

---

## 15. Final Summary

You have successfully completed:

- ✓ Python virtual env setup
- ✓ boto3 & AWS plugin installation
- ✓ AWS credential configuration
- ✓ Dynamic inventory creation
- ✓ Correct hostname mapping ([ip-address](#))
- ✓ SSH authentication fix
- ✓ Successful Ansible ping
- ✓ Running playbooks with dynamic inventory

Your Ansible environment is now fully cloud-ready and scalable.

---