
Observability & Autoscaling in Kubernetes

Observability provides **insight into the health and performance** of your applications and cluster. Autoscaling enables **dynamic resource allocation** based on real-time usage or predefined conditions.

1. Metrics Server

What it is:

A lightweight, in-cluster aggregator of **resource usage data** (CPU and memory) for nodes and pods.

Key Role:

- Provides metrics used by **Horizontal Pod Autoscaler (HPA)** and **kubectl top** commands.

Limitations:

- Only offers recent metrics (no historical data).
- Not suitable for full observability (use Prometheus for that).

Example Usage:

```
kubectl top pod  
kubectl top node
```

2. Horizontal Pod Autoscaler (HPA)

What it does:

Automatically adjusts the number of **replica pods** in a Deployment, StatefulSet, or ReplicaSet **based on CPU/memory usage or custom metrics**.

How it works:

- Queries Metrics Server (or custom metrics).
- Scales pods **up or down** based on threshold (e.g., CPU > 80%).

Example:

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: webapp-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: webapp
  minReplicas: 2
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
      target:
        type: Utilization
        averageUtilization: 75
```

3. Vertical Pod Autoscaler (VPA)

What it does:

Dynamically adjusts the **CPU and memory resource requests/limits** of pods.

When to use:

- For workloads that are not horizontally scalable.

- To avoid over/under-provisioning.

Modes:

- **Auto:** Updates pod specs automatically.
- **Recommendation:** Suggests values but doesn't change anything.
- **Off:** VPA is disabled.

Caveat: Pods are **recreated** to apply new resource settings.



4. Cluster Autoscaler

What it does:

Automatically adjusts the number of **nodes** in your cluster based on unschedulable pods.

How it works:

- Adds nodes if pods can't be scheduled due to resource constraints.
- Removes underutilized nodes to save cost.

Supported platforms:

- AWS, GCP, Azure, and others via cloud provider APIs.

Example Use Case:

Running workloads that spike occasionally — autoscaler can scale up temporarily.



5. Kube-state-metrics

What it is:

A service that exposes **Kubernetes object states** (not resource usage) as Prometheus metrics.

What it exports:

- Deployment status
- Pod readiness
- PVC usage
- CronJob schedules
- Node conditions

Why it's useful:

It provides detailed cluster state metrics for **alerting and dashboarding**.

 6. Event Exporter**What it does:**

Exports **Kubernetes Events** (e.g., failed pod scheduling, node crashes) to external systems like:

- Elasticsearch
- Kafka
- Slack
- Sentry

Why it's useful:

Native K8s events are ephemeral. Exporting them helps with:

- Troubleshooting
 - Alerting
 - Auditing
-

7. Prometheus Operator

What it is:

An Operator that simplifies **deploying and managing Prometheus** instances on Kubernetes.

Key Features:

- CRDs like [ServiceMonitor](#), [PrometheusRule](#) for easy config.
- Automatically discovers and scrapes services.
- Manages Prometheus upgrades, alerts, and rules.

Why use it:

Avoid complex, manual Prometheus setups.

8. Alertmanager

What it does:

Handles **alerts** sent by Prometheus, deduplicates and routes them to:

- Email
- Slack
- PagerDuty

- Webhooks

Features:

- Alert grouping
- Silencing
- Inhibition (prevent alert storms)

Example:

Get notified if CPU usage > 90% across nodes for more than 5 minutes.

9. Grafana

What it is:

A visualization tool that creates **dashboards** for metrics, logs, and traces.

Key Features:

- Integrates with Prometheus, Loki, Jaeger, Elastic, and others.
- Highly customizable dashboards and alerts.

Popular Dashboards:

- Node resource usage
- Pod performance
- HPA activity
- Ingress traffic

10. OpenTelemetry Collector

What it is:

A vendor-neutral component to **collect, process, and export telemetry data** (metrics, logs, traces) from Kubernetes and applications.

Supports:

- Tracing (Jaeger, Zipkin)
- Metrics (Prometheus, OTLP)
- Logs (FluentBit, Loki)


Why it's powerful:


Centralizes telemetry collection in **distributed systems**.

Example Workflow:

App → OpenTelemetry SDK → Otel Collector → Prometheus / Jaeger

How They Work Together (Simplified Flow)





Optional tools:

- Grafana for visualization
- OpenTelemetry Collector for traces/logs
- Cluster Autoscaler & VPA for dynamic resourcing

Summary Table

Tool/Component	Category	Role
Metrics Server	Metrics	Collects pod/node CPU/mem for autoscaling
HPA	Autoscaling	Scales pod replicas based on usage
VPA	Autoscaling	Adjusts pod resource requests/limits
Cluster Autoscaler	Autoscaling	Adds/removes nodes based on demand
Kube-state-metrics	Observability	Exposes Kubernetes object states
Event Exporter	Observability	Forwards events externally
Prometheus Operator	Monitoring	Manages Prometheus setup and configs
Alertmanager	Alerting	Manages alert routing and suppression

Grafana	Visualization	Dashboards for metrics
OpenTelemetry	Observability	Collects traces, logs, and metrics
